

```
In [50]: import pandas as pd
import seaborn as sns
import numpy as np
from matplotlib import pyplot as plt
```

```
In [51]: aerofit = pd.read_csv('Desktop/aerofit_treadmill.txt')
```

```
In [52]: aerofit
```

```
Out[52]:
```

	Product	Age	Gender	Education	MaritalStatus	Usage	Fitness	Income	Miles
0	KP281	18	Male	14	Single	3	4	29562	112
1	KP281	19	Male	15	Single	2	3	31836	75
2	KP281	19	Female	14	Partnered	4	3	30699	66
3	KP281	19	Male	12	Single	3	3	32973	85
4	KP281	20	Male	13	Partnered	4	2	35247	47
...
175	KP781	40	Male	21	Single	6	5	83416	200
176	KP781	42	Male	18	Single	5	4	89641	200
177	KP781	45	Male	16	Single	5	5	90886	160
178	KP781	47	Male	18	Partnered	4	5	104581	120
179	KP781	48	Male	18	Partnered	4	5	95508	180

180 rows × 9 columns

```
In [53]: aerofit.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 180 entries, 0 to 179
Data columns (total 9 columns):
#   Column          Non-Null Count  Dtype
---  -
0   Product         180 non-null   object
1   Age             180 non-null   int64
2   Gender          180 non-null   object
3   Education        180 non-null   int64
4   MaritalStatus    180 non-null   object
5   Usage           180 non-null   int64
6   Fitness         180 non-null   int64
7   Income          180 non-null   int64
8   Miles           180 non-null   int64
dtypes: int64(6), object(3)
memory usage: 10.6+ KB
```

In [54]: aerofit.describe()

Out[54]:

	Age	Education	Usage	Fitness	Income	Miles
count	180.000000	180.000000	180.000000	180.000000	180.000000	180.000000
mean	28.788889	15.572222	3.455556	3.311111	53719.577778	103.194444
std	6.943498	1.617055	1.084797	0.958869	16506.684226	51.863605
min	18.000000	12.000000	2.000000	1.000000	29562.000000	21.000000
25%	24.000000	14.000000	3.000000	3.000000	44058.750000	66.000000
50%	26.000000	16.000000	3.000000	3.000000	50596.500000	94.000000
75%	33.000000	16.000000	4.000000	4.000000	58668.000000	114.750000
max	50.000000	21.000000	7.000000	5.000000	104581.000000	360.000000

In [55]: aerofit.notna().sum()


Out[55]: Product 180
 Age 180
 Gender 180
 Education 180
 MaritalStatus 180
 Usage 180
 Fitness 180
 Income 180
 Miles 180
 dtype: int64

In [56]: aerofit.nunique()

Out[56]: Product 3
 Age 32
 Gender 2
 Education 8
 MaritalStatus 2
 Usage 6
 Fitness 5
 Income 62
 Miles 37
 dtype: int64

```
In [57]: sns.scatterplot(data=aerofit,x='Age',y='Income',hue='Product')
```


```
Out[57]: <AxesSubplot:xlabel='Age', ylabel='Income'>
```



```
In [58]: #KP281 is mostly bought by people with income>50000  
#KP481 and 781 is mostly bought by people with income<50000
```


```
In [59]: sns.countplot(data=aerofit,x='Gender',hue='Product')  
#Females have not bought 781, for it target base should be male
```

```
Out[59]: <AxesSubplot:xlabel='Gender', ylabel='count'>
```



```
In [60]: sns.countplot(data=aerofit,x='Education',hue='Product')
```

```
Out[60]: <AxesSubplot:xlabel='Education', ylabel='count'>
```




```
In [134]: print(aerofit.value_counts('Product'))
print(aerofit.value_counts('Gender'))
print(aerofit.value_counts('MaritalStatus'))
print(aerofit.value_counts('Usage'))
print(aerofit.value_counts('Fitness'))
#most sevalue_countsct is 281
```

```
Product
KP281    80
KP481    60
KP781    40
dtype: int64
Gender
Male     104
Female   76
dtype: int64
MaritalStatus
Partnered 107
Single    73
dtype: int64
Usage
3     69
4     52
2     33
5     17
6      7
7      2
dtype: int64
Fitness
3     97
5     31
2     26
4     24
1      2
dtype: int64
```


```
In [62]: sns.countplot(data=aerofit,x='Fitness',hue='Product')
```

```
Out[62]: <AxesSubplot:xlabel='Fitness', ylabel='count'>
```




```
In [63]: sns.countplot(data=aerofit,x='MaritalStatus',hue='Product')
```

```
Out[63]: <AxesSubplot:xlabel='MaritalStatus', ylabel='count'>
```



```
In [64]: sns.histplot(data=aerofit,x='Age',hue='Product')
```

```
Out[64]: <AxesSubplot:xlabel='Age', ylabel='Count'>
```




```
In [65]: pd.crosstab(columns=aerofit['Product'],index=1)
```

```
Out[65]:
```

	Product	KP281	KP481	KP781
row_0	1	80	60	40

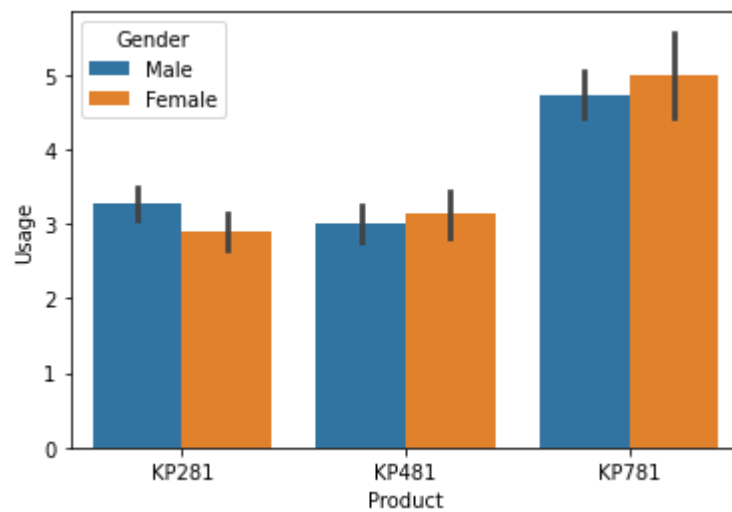
```
In [66]: sns.barplot(data=aerofit,x='Product',y='Miles')
```

```
Out[66]: <AxesSubplot:xlabel='Product', ylabel='Miles'>
```



```
In [67]: sns.barplot(data=aerofit,x='Product',y='Usage',hue='Gender')
```

```
Out[67]: <AxesSubplot:xlabel='Product', ylabel='Usage'>
```



```
In [68]: sns.pairplot(data=aerofit)
```

```
Out[68]: <seaborn.axisgrid.PairGrid at 0x14cf2cd0>
```



```
In [ ]:
```


In [69]: *#two-way contingency tables*

```
pd.crosstab([aerofit.Age, aerofit.Gender], aerofit.Product, margins = False)
```

Out[69]:

		Product		
		KP281	KP481	KP781
Age	Gender			
18	Male	1	0	0
19	Female	1	0	0
	Male	2	1	0
20	Female	1	1	0
	Male	1	2	0
21	Female	2	1	0
	Male	2	2	0
22	Female	3	0	0
	Male	1	0	3
23	Female	3	3	1
	Male	5	4	2
24	Female	3	2	1
	Male	2	1	3
25	Female	4	5	1
	Male	3	6	6
26	Female	3	2	1
	Male	4	1	1
27	Female	2	0	0
	Male	1	1	3
28	Female	4	0	1
	Male	2	0	2
29	Female	2	1	0
	Male	1	0	2
30	Female	0	2	1
	Male	2	0	2
31	Female	1	2	0
	Male	1	1	1
32	Female	1	0	0
	Male	1	2	0
33	Female	2	3	1
	Male	0	2	0

	Product	KP281	KP481	KP781
Age	Gender			
34	Female	1	1	0
	Male	1	2	1
35	Female	2	2	0
	Male	1	2	1
36	Male	1	0	0
37	Female	1	1	0
38	Female	1	1	0
	Male	3	1	1
39	Male	1	0	0
40	Female	0	2	0
	Male	1	1	1
41	Male	1	0	0
42	Male	0	0	1
43	Male	1	0	0
44	Female	1	0	0
45	Male	0	1	1
46	Female	1	0	0
47	Male	1	0	1
48	Male	0	1	1
50	Female	1	0	0

```
In [70]: aerofit.isnull().sum()
#No null values
```

```
Out[70]: Product      0
Age      0
Gender    0
Education 0
MaritalStatus 0
Usage     0
Fitness   0
Income    0
Miles     0
dtype: int64
```

```
In [81]: prob1 = pd.crosstab([aerofit.MaritalStatus], aerofit.Product, margins = False)
prob1
```

Out[81]:

	Product	KP281	KP481	KP781
MaritalStatus				
Partnered		48	36	23
Single		32	24	17

```
In [93]: #probability
print(prob1.loc[['Single']][0]/prob1.loc['Single'].sum())
print(prob1.loc[['Partnered']][0]/prob1.loc['Partnered'].sum())
```

```
Product
KP281    0.438356
KP481    0.328767
KP781    0.232877
Name: Single, dtype: float64
Product
KP281    0.448598
KP481    0.336449
KP781    0.214953
Name: Partnered, dtype: float64
```

```
In [97]: prob2=pd.crosstab([aerofit.Usage, aerofit.Fitness], aerofit.Product, margins = False)
prob2
```

Out[97]:

		Product		
		KP281	KP481	KP781
Usage	Fitness			
2	1	0	1	0
	2	9	5	0
	3	10	8	0
3	1	1	0	0
	2	4	6	0
	3	27	20	0
	4	5	5	0
	5	0	0	1
4	2	1	1	0
	3	17	9	4
	4	3	2	2
	5	1	0	12
5	3	0	2	0
	4	1	1	4
	5	1	0	8
6	4	0	0	1
	5	0	0	6
7	5	0	0	2

```
In [99]: #probability
print(prob2.loc[[2][0]]/prob2.loc[2].sum())
print(prob2.loc[[3][0]]/prob2.loc[3].sum())
print(prob2.loc[[4][0]]/prob2.loc[4].sum())
print(prob2.loc[[5][0]]/prob2.loc[5].sum())
print(prob2.loc[[6][0]]/prob2.loc[6].sum())
print(prob2.loc[[7][0]]/prob2.loc[7].sum())
```

Product	KP281	KP481	KP781
Fitness			
1	0.000000	0.071429	NaN
2	0.473684	0.357143	NaN
3	0.526316	0.571429	NaN

Product	KP281	KP481	KP781
Fitness			
1	0.027027	0.000000	0.0
2	0.108108	0.193548	0.0
3	0.729730	0.645161	0.0
4	0.135135	0.161290	0.0
5	0.000000	0.000000	1.0

Product	KP281	KP481	KP781
Fitness			
2	0.045455	0.083333	0.000000
3	0.772727	0.750000	0.222222
4	0.136364	0.166667	0.111111
5	0.045455	0.000000	0.666667

Product	KP281	KP481	KP781
Fitness			
3	0.0	0.666667	0.000000
4	0.5	0.333333	0.333333
5	0.5	0.000000	0.666667

Product	KP281	KP481	KP781
Fitness			
4	NaN	NaN	0.142857
5	NaN	NaN	0.857143

Product	KP281	KP481	KP781
Fitness			
5	NaN	NaN	1.0

```
In [103]: prob3=pd.crosstab([aerofit.Education], aerofit.Product, margins = False)
          prob3
```

Out[103]:

Product	KP281	KP481	KP781
Education			
12	2	1	0
13	3	2	0
14	30	23	2
15	4	1	0
16	39	31	15
18	2	2	19
20	0	0	1
21	0	0	3

```
In [106]: print(prob3.loc[[12][0]]/prob3.loc[12].sum())
print(prob3.loc[[13][0]]/prob3.loc[13].sum())
print(prob3.loc[[14][0]]/prob3.loc[14].sum())
print(prob3.loc[[15][0]]/prob3.loc[15].sum())
print(prob3.loc[[16][0]]/prob3.loc[16].sum())
print(prob3.loc[[18][0]]/prob3.loc[18].sum())
print(prob3.loc[[20][0]]/prob3.loc[20].sum())
print(prob3.loc[[21][0]]/prob3.loc[21].sum())
```

Product

KP281 0.666667

KP481 0.333333

KP781 0.000000

Name: 12, dtype: float64

Product

KP281 0.6

KP481 0.4

KP781 0.0

Name: 13, dtype: float64

Product

KP281 0.545455

KP481 0.418182

KP781 0.036364

Name: 14, dtype: float64

Product

KP281 0.8

KP481 0.2

KP781 0.0

Name: 15, dtype: float64

Product

KP281 0.458824

KP481 0.364706

KP781 0.176471

Name: 16, dtype: float64

Product

KP281 0.086957

KP481 0.086957

KP781 0.826087

Name: 18, dtype: float64

Product

KP281 0.0

KP481 0.0

KP781 1.0

Name: 20, dtype: float64

Product

KP281 0.0

KP481 0.0

KP781 1.0

Name: 21, dtype: float64

```
In [108]: prob4=pd.crosstab([aerofit.Gender], aerofit.Product, margins = False)
prob4
```

```
Out[108]:
```

Product	KP281	KP481	KP781
Gender			
Female	40	29	7
Male	40	31	33

```
In [124]: print(prob4.loc[['Female']][0]/prob4.loc['Female'].sum())
print(prob4.loc[['Male']][0]/prob4.loc['Male'].sum())
```

```
Product
KP281    0.526316
KP481    0.381579
KP781    0.092105
Name: Female, dtype: float64
Product
KP281    0.384615
KP481    0.298077
KP781    0.317308
Name: Male, dtype: float64
```

```
In [73]: #convert categorical into numerical values
```

```
In [74]: #Creating new df
aerofit1=aerofit
```

```
In [75]: aerofit1
```

```
Out[75]:
```

	Product	Age	Gender	Education	MaritalStatus	Usage	Fitness	Income	Miles
0	KP281	18	Male	14	Single	3	4	29562	112
1	KP281	19	Male	15	Single	2	3	31836	75
2	KP281	19	Female	14	Partnered	4	3	30699	66
3	KP281	19	Male	12	Single	3	3	32973	85
4	KP281	20	Male	13	Partnered	4	2	35247	47
...
175	KP781	40	Male	21	Single	6	5	83416	200
176	KP781	42	Male	18	Single	5	4	89641	200
177	KP781	45	Male	16	Single	5	5	90886	160
178	KP781	47	Male	18	Partnered	4	5	104581	120
179	KP781	48	Male	18	Partnered	4	5	95508	180

180 rows × 9 columns


```
In [76]: #aerofit1['Product']=aerofit1['Product'].map({'KP281':0, 'KP481':1, 'KP781':2})
#aerofit1['Gender']=aerofit1['Gender'].map({'Male':0, 'Female':1})
#aerofit1['MaritalStatus']=aerofit1['MaritalStatus'].map({'Single':0, 'Partnered':1})
```

```
In [77]: aerofit1
```


```
Out[77]:
```

	Product	Age	Gender	Education	MaritalStatus	Usage	Fitness	Income	Miles
0	KP281	18	Male	14	Single	3	4	29562	112
1	KP281	19	Male	15	Single	2	3	31836	75
2	KP281	19	Female	14	Partnered	4	3	30699	66
3	KP281	19	Male	12	Single	3	3	32973	85
4	KP281	20	Male	13	Partnered	4	2	35247	47
...
175	KP781	40	Male	21	Single	6	5	83416	200
176	KP781	42	Male	18	Single	5	4	89641	200
177	KP781	45	Male	16	Single	5	5	90886	160
178	KP781	47	Male	18	Partnered	4	5	104581	120
179	KP781	48	Male	18	Partnered	4	5	95508	180

180 rows × 9 columns

```
In [78]: sns.heatmap(aerofit1.corr(), cmap="YlGnBu", annot=True)
```

```
Out[78]: <AxesSubplot:>
```



In [79]:

aerofit1.corr()

Out[79]:

	Age	Education	Usage	Fitness	Income	Miles
Age	1.000000	0.280496	0.015064	0.061105	0.513414	0.036618
Education	0.280496	1.000000	0.395155	0.410581	0.625827	0.307284
Usage	0.015064	0.395155	1.000000	0.668606	0.519537	0.759130
Fitness	0.061105	0.410581	0.668606	1.000000	0.535005	0.785702
Income	0.513414	0.625827	0.519537	0.535005	1.000000	0.543473
Miles	0.036618	0.307284	0.759130	0.785702	0.543473	1.000000

In []:

In []:

In []:

In []:

In []:

In []:

In []:

In []:

In []: